

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import classification_report, confusion_matrix, _
    accuracy_score

import warnings
warnings.filterwarnings("ignore")
```

```
[2]: # Load the data
df = pd.read_csv('/content/heart.csv')
df.head()
```

```
[2]: Age Sex ChestPainType RestingBP Cholesterol FastingBS RestingECG MaxHR
\
```

```
0  40M    ATA  140  289  0    Normal    172
1  49 F  NAP  160  180  0  Normal  156  2  37 M  ATA  130  283  0  ST  98
3  48F    ASY  138  214  0    Normal    108
4  54M    NAP  150  195  0    Normal    122
```

```
ExerciseAngina Oldpeak ST_Slope HeartDisease
0              N  0.0    Up    0
1              N  1.0  Flat  1
2              N  0.0    Up    0
3              Y  1.5  Flat  1
4              N  0.0    Up    0
```

```
[3]: df.info()
```

```
<class
'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to
917 Data columns (total 12
columns):
#   Column          Non-Null Count  Dtype
---  -
0  Age            918 non-null    int64
1  Sex            918 non-null    object
```

```

2  ChestPainType      918 non-null    object
3  RestingBP         918 non-null    int64
4  Cholesterol       918 non-null    int64
5  FastingBS        918 non-null    int64
6  RestingECG       918 non-null    object
7  MaxHR            918 non-null    int64
8  ExerciseAngina    918 non-null    object
9  Oldpeak          918 non-null    float64
10 ST_Slope         918 non-null    object
11 HeartDisease      918 non-null    int64

```

```

dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB

```

```
[4]: df.describe()
```

```

[4]:      Age  RestingBP  Cholesterol  FastingBS  MaxHR \
count  918.000000  918.000000      918.000000  918.000000
mean    53.510893  132.396514  198.799564    0.233115
std      9.432617   18.514154  109.384145    0.423046
min     28.000000    0.000000    0.000000    0.000000
25%     47.000000  120.000000  173.250000    0.000000
50%     54.000000  130.000000  223.000000    0.000000
75%     60.000000  140.000000  267.000000    0.000000
max     77.000000  200.000000  603.000000    1.000000

      Oldpeak  HeartDisease
count  918.000000  918.000000
mean    0.887364  0.553377
std     1.066570  0.497414
min     0.000000  0.000000
25%     0.000000  0.000000
50%     0.600000  1.000000
75%     1.500000  1.000000
max     6.200000  1.000000

```

```
[5]: df.isnull().sum()
```

```
[5]: Age      0
```

```

Sex          0
ChestPainType 0
RestingBP    0
Cholesterol  0
FastingBS    0
RestingECG   0
MaxHR        0
ExerciseAngina 0
Oldpeak      0
ST_Slope     0
HeartDisease 0
dtype: int64

```

```

[6]: # Check for missing values (already confirmed none in previous
      # step) # Handle any zero values that might be errors (like
      Cholesterol=0) df['Cholesterol'] = df['Cholesterol'].replace(0,
      df['Cholesterol'].median())

```

```

[7]: # Check for duplicates print(f"Number of
      duplicates: {df.duplicated().sum()}") df =
      df.drop_duplicates()

```

```

Number of duplicates: 0

```

```

[8]: # Check target variable distribution
      df['HeartDisease'].value_counts(normalize=True)

```

```

[8]: HeartDisease
1    0.553377
0    0.446623
Name: proportion, dtype: float64

```

```

[9]: # Separate features and
      target X =
      df.drop('HeartDisease',
      axis=1) y = df['HeartDisease']

      # Identify categorical and numerical columns categorical_cols
      = ['Sex', 'ChestPainType', 'FastingBS', 'RestingECG',
      'ExerciseAngina', 'ST_Slope'] numerical_cols = ['Age',
      'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak']

      # Create preprocessing pipelines
      numerical_transformer = StandardScaler()
      categorical_transformer = OneHotEncoder(handle_unknown='ignore')

      preprocessor = ColumnTransformer(
          transformers=[

```

```
        ('num', numerical_transformer, numerical_cols),
        ('cat', categorical_transformer,
 categorical_cols) ])

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2,
↳random_state=42, stratify=y)
```

```

[10]: from sklearn.svm import SVC
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.neighbors import KNeighborsClassifier
      from xgboost import XGBClassifier

      # Initialize models
      models = {
          'SVM': SVC(random_state=42),
          'Decision Tree': DecisionTreeClassifier(random_state=42),
          'Random Forest': RandomForestClassifier(random_state=42),
          'KNN': KNeighborsClassifier(),
          'XGBoost': XGBClassifier(random_state=42, eval_metric='logloss')
      }

      # Create a dictionary to store results
      results = {}

      # Train and evaluate each model
      for name, model in models.items():
          # Create pipeline
          pipeline = Pipeline(steps=[
              ('preprocessor', preprocessor),
              ('classifier', model)
          ])

          # Train model
          pipeline.fit(X_train, y_train)

          # Make predictions
          y_pred = pipeline.predict(X_test)

          # Evaluate model
          accuracy = accuracy_score(y_test, y_pred)
          report = classification_report(y_test, y_pred)
          cm = confusion_matrix(y_test, y_pred)

          # Store results
          results[name] = {
              'model': pipeline,
              'accuracy': accuracy,
              'report': report,
              'confusion_matrix': cm
          }

      print(f"\n{name} Results:")

```

SVM Results:

Decision Tree Results:

Random Forest Results:

KNN Results:

XGBoost Results:

```
[11]: print(f"Accuracy: {accuracy:.4f}")
```

Accuracy: 0.8370

```
[12]: print("Classification Report:")
      print(report)
```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.84	0.82	82
1	0.87	0.83	0.85	102
accuracy			0.84	184
macro avg	0.83	0.84	0.84	184
weighted avg	0.84	0.84	0.84	184

```
[13]: print("Confusion Matrix:")
      print(cm)
```

Confusion Matrix:

```
[[69 13]
 [17 85]]
```

```
[14]: # Compare model accuracies accuracies = {name: result['accuracy']
      for name, result in results.items()} sorted_accuracies =
      sorted(accuracies.items(), key=lambda x: x[1], reverse=True)

      print("\nModel Accuracy
      Comparison:") for name, acc in
      sorted_accuracies: print(f"{name}:
      {acc:.4f}")

      # Visualize accuracy comparison
      plt.figure(figsize=(10, 6))
      plt.bar(accuracies.keys(),
      accuracies.values()) plt.title('Model Accuracy
      Comparison')
```

```

plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.ylim(0.7, 1.0)
plt.xticks(rotation=45)
plt.show()

# Feature importance for tree-based models
for name in ['Decision Tree', 'Random Forest', 'XGBoost']:
    try:
        # Get feature names after one-hot encoding
        preprocessor.fit(X_train)
        feature_names = numerical_cols + list(preprocessor.
        ↪named_transformers_['cat'].get_feature_names_out(categorical_cols))

        # Get feature importances
        if name == 'XGBoost':
            importances = results[name]['model'].named_steps['classifier'].
            ↪feature_importances_
        else:
            importances = results[name]['model'].named_steps['classifier'].
            ↪feature_importances_

        # Create DataFrame for visualization
        importance_df = pd.DataFrame({'Feature': feature_names, 'Importance': ↪
        ↪importances})
        importance_df = importance_df.sort_values('Importance', ↪
        ↪ascending=False).head(10)

        # Plot
        plt.figure(figsize=(10, 6))
        plt.title(f'{name} - Top 10 Feature Importances')
        sns.barplot(x='Importance', y='Feature', data=importance_df)
        plt.show()
    except Exception as e:
        print(f"Could not plot feature importance for {name}: {str(e)}")

```

Model Accuracy Comparison:



KNN: 0.8859


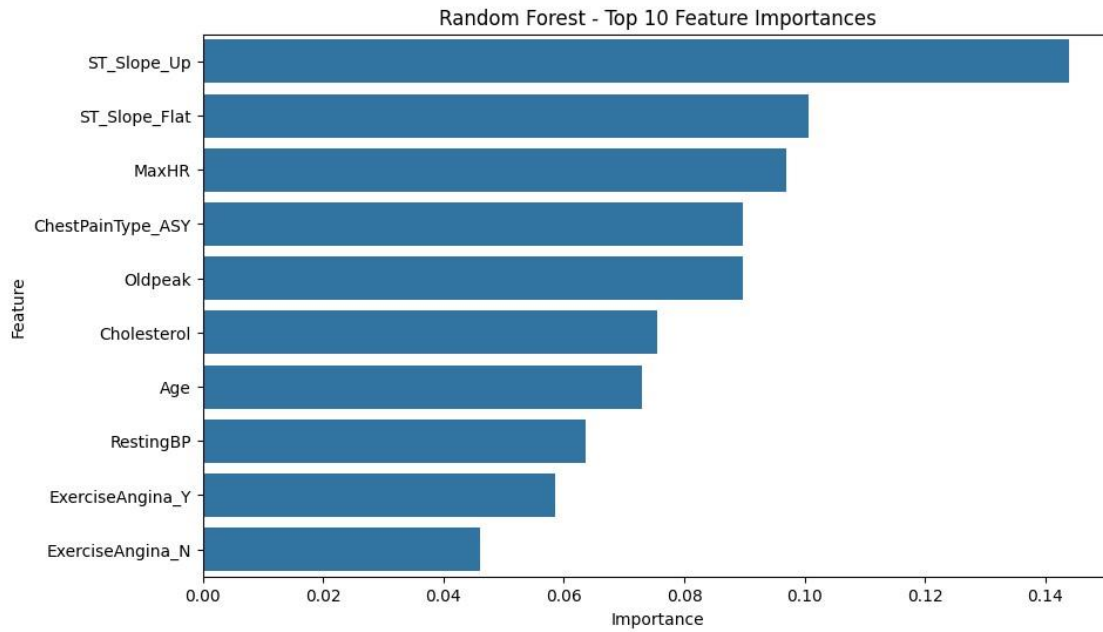
SVM: 0.8804

Random Forest: 0.8750

XGBoost: 0.8370

Decision Tree: 0.7609





```
[15]: from sklearn.model_selection import GridSearchCV

# Example for Random Forest
param_grid = {
    'classifier__n_estimators': [100, 200, 300],
```

```

        'classifier__max_depth': None, 5, 10],
        'classifier__min_samples_split': 2, 5, 10]
    }

    rf_pipeline = Pipeline(steps=[
        ('preprocessor', preprocessor),
        ('classifier', RandomForestClassifier(random_state=42))
    ])

    grid_search = GridSearchCV(rf_pipeline, param_grid, cv=5, scoring='accuracy',
                               n_jobs=-1)
    grid_search.fit(X_train, y_train)

    print(f"Best parameters: {grid_search.best_params_}")
    print(f"Best accuracy: {grid_search.best_score_:.4f}")

    # Update the best model in results
    results['Random Forest (Tuned)'] = {
        'model': grid_search.best_estimator_,
        'accuracy': accuracy_score(y_test, grid_search.best_estimator_.
        predict(X_test)),
        'report': classification_report(y_test, grid_search.best_estimator_.
        predict(X_test)),
        'confusion_matrix': confusion_matrix(y_test, grid_search.best_estimator_.
        predict(X_test))
    }

```

```

Best parameters: {'classifier__max_depth': 5,
'classifier__min_samples_split':
10, 'classifier__n_estimators':
200} Best accuracy: 0.8637

```

```

[16]: from sklearn.neighbors import KNeighborsClassifier from
sklearn.model_selection import GridSearchCV,
RandomizedSearchCV from sklearn.metrics import
make_scorer, accuracy_score, f1_score import numpy as np

```

```

[17]: param_grid = {
        'classifier__n_neighbors': np.arange(3, 30, 2), # Odd numbers to
        avoid ties
        'classifier__weights': ['uniform', 'distance'],
        'classifier__p': [1, 2], # 1: Manhattan, 2: Euclidean
        'classifier__metric': ['minkowski', 'cosine',
        'manhattan'] }

```

```
[18]: # Reuse the preprocessor from previous steps
knn_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', KNeighborsClassifier())
])
```

```
[19]: # Using accuracy as the scoring metric
scorer = make_scorer(accuracy_score)

# GridSearchCV with 5-fold cross-validation
grid_search = GridSearchCV(
    estimator=knn_pipeline,
    param_grid=param_grid,
    scoring=scorer,
    cv=5,
    n_jobs=-1, # Use all available CPU cores
    verbose=1
)
```

```
[20]: grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 168 candidates, totalling 840 fits

```
[20]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('preprocessor',
                                             ColumnTransformer(transformers=[('num',
                                                                                   StandardScaler(),
                                                                                   ['Age',
                                                                                   'RestingBP',
                                                                                   'Cholesterol',
                                                                                   'MaxHR',
                                                                                   'Oldpeak'])],
                                             ('cat',
                                             OneHotEncoder(handle_unknown='ignore'),
                                             ['Sex',
                                             'ChestPainType',
                                             'FastingBS',
                                             'RestingECG',
                                             'ExerciseAngina',
                                             'ST_Slope'])])),
                  ('classifier', KNeighborsClassifier())]),
                  n_jobs=-1,
                  param_grid={'classifier__metric': ['minkowski', 'cosine',
                                                       'manhattan'],
                              'classifier__n_neighbors': array([ 3, 5, 7, 9, 11,
```

```

13, 15, 17, 19, 21, 23, 25, 27, 29]),
        'classifier__p': [1,
                           2],
        'classifier__weights': ['uniform',
                                'distance']},
        scoring=make_scorer(accuracy_score,
                             response_method='predict'), verbose=1)
[21]: # Best parameters and score print("Best parameters found: ",
grid_search.best_params_) print("Best cross-validation accuracy:
{:.2f}%".format(grid_search.best_score_ * 100))

# Evaluate on test set
best_knn = grid_search.best_estimator_ y_pred =
best_knn.predict(X_test) test_accuracy = accuracy_score(y_test,
y_pred) print("\nTest set accuracy with best KNN:
{:.2f}%".format(test_accuracy * 100))

```

```

Best parameters found: {'classifier__metric': 'minkowski',
'classifier__n_neighbors': np.int64(15), 'classifier__p': 1,
'classifier__weights': 'uniform'}
Best cross-validation accuracy:

```

```

86.37% Test set accuracy with best

```

```

KNN: 91.30%

```

```

[22]: # Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

```

Classification Report:

```

	precision	recall	f1-score	support
0	0.90	0.90	0.90	82
1	0.92	0.92	0.92	102
accuracy			0.91	184
macro avg	0.91	0.91	0.91	184
weighted	0.91	0.91	0.91	184
avg				

```

[23]: # Confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

Confusion Matrix:

```
[[74 8]
 [ 8 94]]
```

```
[24]: # Get the best estimator from GridSearchCV
best_knn = grid_search.best_estimator_

# View the best parameters
print("Best Parameters Found:")

print(grid_search.best_params_)
```

Best Parameters Found:

```
{'classifier__metric': 'minkowski', 'classifier__n_neighbors':
np.int64(15),
'classifier__p': 1, 'classifier__weights': 'uniform'}
```

```
[25]: # Get the best estimator from GridSearchCV
best_knn = grid_search.best_estimator_

# View the best parameters
print("Best Parameters Found:")
print(grid_search.best_params_)
```

Best Parameters Found:

```
{'classifier__metric': 'minkowski', 'classifier__n_neighbors':
np.int64(15),
'classifier__p': 1, 'classifier__weights': 'uniform'}
```

```
[26]: # Predict on test set
y_pred = best_knn.predict(X_test)
y_pred_proba = best_knn.predict_proba(X_test)[:, 1] # Probability estimates
↳ for class 1
```

```
[27]: from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)
print(f"\nTest Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")
```

Test Accuracy: 0.9130 (91.30%)

```
[28]: from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)
print(f"\nTest Accuracy: {accuracy:.4f}
({accuracy*100:.2f}%)")
```

Test Accuracy: 0.9130 (91.30%)

```
[29]: from sklearn.metrics import classification_report

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.90	0.90	0.90	82
1	0.92	0.92	0.92	102
accuracy			0.91	184
macro avg	0.91	0.91	0.91	184
weighted avg	0.91	0.91	0.91	184


```
[30]: from sklearn.metrics import confusion_matrix
import seaborn as sns

cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(cm)

# Visualize confusion matrix
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Heart Disease', 'Heart Disease'],
            yticklabels=['No Heart Disease', 'Heart Disease'])
plt.title('Confusion Matrix')
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show()
```

Confusion Matrix:


```
[[74  8]
 [ 8 94]]
```



```
[31]: from sklearn.metrics import roc_curve, roc_auc_score

fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
auc_score = roc_auc_score(y_test, y_pred_proba)

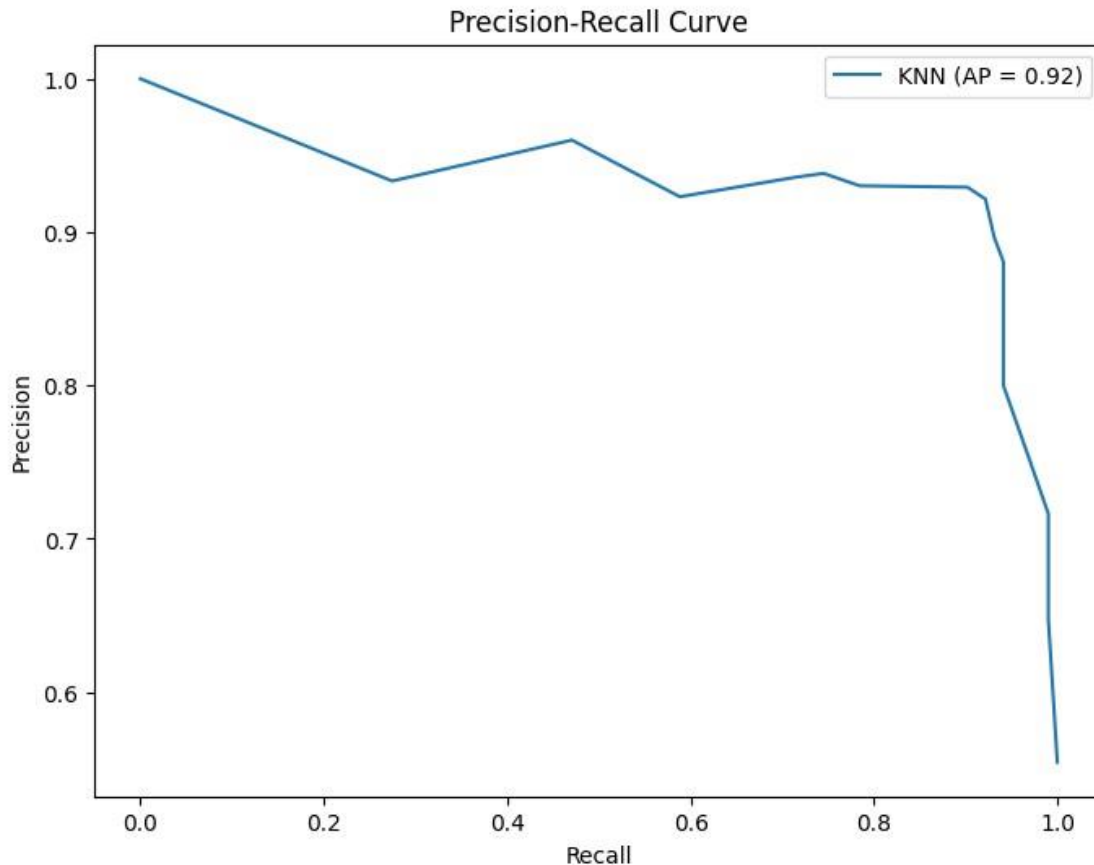
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'KNN (AUC = {auc_score:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```



```
[32]: from sklearn.metrics import precision_recall_curve,
average_precision_score

precision, recall, _ = precision_recall_curve(y_test, y_pred_proba)
avg_precision = average_precision_score(y_test, y_pred_proba)

plt.figure(figsize=(8, 6)) plt.plot(recall, precision,
label=f'KNN (AP = {avg_precision:.2f})')
plt.xlabel('Recall') plt.ylabel('Precision')
plt.title('Precision-Recall Curve') plt.legend() plt.show()
```

```
[33]: from sklearn.inspection import permutation_importance


# Get feature names after preprocessing
preprocessor.fit(X_train) feature_names = numerical_cols +
list(preprocessor.named_transformers_['cat'].
get_feature_names_out(categorical_cols))

# Calculate permutation importance result =
permutation_importance(best_knn, X_test, y_test,
n_repeats=10, random_state=42)

# Sort features by importance
sorted_idx = result.importances_mean.argsort()[::-1]

# Plot top 10 features
plt.figure(figsize=(10, 6))
plt.barh(np.array(feature_names)[sorted_idx][:10],
result.importances_mean[sorted_idx][:10])
plt.xlabel("Permutation Importance")
```

```
plt.title("Top 10 Important Features")
plt.gca().invert_yaxis()
plt.show()
```



```
[34]: from sklearn.metrics import precision_score, recall_score, f1_score

print("\nFinal Evaluation Metrics:")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(f"Precision: {precision_score(y_test, y_pred):.4f}")
print(f"Recall: {recall_score(y_test, y_pred):.4f}")
print(f"F1 Score: {f1_score(y_test, y_pred):.4f}")
print(f"ROC AUC: {roc_auc_score(y_test, y_pred_proba):.4f}")
print(f"Average Precision: {avg_precision: .4f}")
```

```
Final Evaluation Metrics:
Accuracy: 0.9130
Precision: 0.9216
Recall: 0.9216
F1 Score: 0.9216
ROC AUC: 0.9338
Average Precision: 0.9228
```

```
[35]: import pandas as pd  
import numpy as np
```

```

def predict_heart_disease(model, input_data):
    """
    Predicts the probability of heart disease using the trained KNN model.

    Parameters:
    -----
    model : Pipeline
        The trained scikit-learn pipeline (including preprocessing and
        ↪ classifier)
    input_data : dict or pandas.DataFrame
        Input features for prediction. Can be:
        - A dictionary of feature names and values
        - A pandas DataFrame with one row of data

    Returns:
    -----
    dict
        A dictionary containing:
        - 'prediction': 0 (no heart disease) or 1 (heart disease)
        - 'probability': Probability of heart disease (0-1)
        - 'interpretation': Text description of the result
    """

    # Define expected features and their validation ranges
    feature_ranges = {
        'Age': (20, 100),
        'Sex': ['M', 'F'],
        'ChestPainType': ['ATA', 'NAP', 'ASY', 'TA'],
        'RestingBP': (80, 200),
        'Cholesterol': (100, 600),
        'FastingBS': [0, 1],
        'RestingECG': ['Normal', 'ST', 'LVH'],
        'MaxHR': (60, 220),
        'ExerciseAngina': ['Y', 'N'],
        'Oldpeak': (-2.5, 6.5),
        'ST_Slope': ['Up', 'Flat', 'Down']
    }

    try:
        # Convert input to DataFrame if it's a dictionary
        if isinstance(input_data, dict):
            input_df = pd.DataFrame([input_data])
        else:
            input_df = input_data.copy()

        # Validate input features
        missing_features = set(feature_ranges.keys()) - set(input_df.columns)

```

```

if missing_features:
    raise ValueError(f"Missing features: {missing_features}")

# Validate feature values
for feature, valid_range in feature_ranges.items():
    value = input_df[feature].iloc[0]

    if feature in ['Sex', 'ChestPainType', 'FastingBS', 'RestingECG', 'ExerciseAngina', 'ST_Slope']:
        if value not in valid_range:
            raise ValueError(f"Invalid value for {feature}. Must be one of: {valid_range}")
        else:
            if not (valid_range[0] <= value <= valid_range[1]):
                raise ValueError(f"Invalid value for {feature}. Must be between {valid_range[0]} and {valid_range[1]}")

# Make prediction
proba = model.predict_proba(input_df)[0][1]
prediction = model.predict(input_df)[0]

# Create interpretation
if prediction == 1:
    interpretation = f"High risk of heart disease ({proba*100:.1f}% probability)"
else:
    interpretation = f"Low risk of heart disease ({(1-proba)*100:.1f}% probability)"

return {
    'prediction': int(prediction),
    'probability': float(proba),
    'interpretation': interpretation
}

except Exception as e:
    return {
        'error': str(e),
        'suggestion': 'Please check your input data format and values'
    }

```

```

[36]: import joblib

# Save the model
joblib.dump(best_knn, 'heart_disease_knn_model.pkl')

# Later, load the model

```

```
model = joblib.load('heart_disease_knn_model.pkl')
```

```
[37]: # Example input (as dictionary)
```

```
patient_data = {
    'Age': 52,
    'Sex': 'M',
    'ChestPainType': 'ASY',
    'RestingBP': 125,
    'Cholesterol': 212,
    'FastingBS': 0,
    'RestingECG': 'Normal',
    'MaxHR': 168,
    'ExerciseAngina': 'N',
    'Oldpeak': 1.0,
    'ST_Slope': 'Flat'
}

# Get prediction
result = predict_heart_disease(model, patient_data)
print(result)
```

```
{'prediction': 1, 'probability': 0.7333333333333333,
 'interpretation': 'High risk of heart disease (73.3% probability)'}
```

```
[38]: from sklearn.model_selection import RandomizedSearchCV
```

```
random_search = RandomizedSearchCV(
    estimator=knn_pipeline,
    param_distributions=param_grid,
    n_iter=50, # Number of parameter settings sampled
    scoring=scorer,
    cv=5,
    n_jobs=-1,
    verbose=1,
    random_state=42
)

random_search.fit(X_train, y_train)

# Analyze results similarly to grid search
```

Fitting 5 folds for each of 50 candidates, totalling 250 fits

```
[38]: RandomizedSearchCV(cv=5,
                        estimator=Pipeline(steps=[('preprocessor',
                                                    ColumnTransformer(transformers=[('num',
                                                                                          StandardScaler()),
```

```

['Age',
 'RestingBP',
 'Cholesterol',
 'MaxHR',
 'Oldpeak']],
('cat',
OneHotEncoder(handle_unknown='ignore'),
['Sex',
 'ChestPainType',
 'FastingBS',
 'RestingECG',
 'ExerciseAngina',
 'ST_Slope']]])),

        ('classifier',
         KNeighborsClassifier()))],
n_iter=50, n_jobs=-1,
param_distributions={'classifier__metric':
 ['minkowski',
                                'cosine',
                                'manhattan'],
                    'classifier__n_neighbors':
array([ 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29]),
                    'classifier__p': [1, 2],
                    'classifier__weights': ['uniform',
                                             'distance']},

        random_state=42,
        scoring=make_scorer(accuracy_score,
response_method='predict'),
        verbose=1)

```